

Adaptive Methods for Nonconvex Optimization

Grant Block, Duke Kwon, Priya Sapro

Authors and Publication Date

- Authors
 - Shahnk K. Reddi
 - Manzil Zaheer
 - Devendra Sachan
 - Satyen Kale
 - Sanjiv Kumar
- Published in NeurIPS (Conference on Neural Information Processing Systems), 2018

The Flaws of Current First Order EMA Methods & Motivation

- Exponential moving Average (EMA) methods like RMSProp and Adam fail to converge in certain convex settings
 - Quickly forget gradient information
 - Current gradient is not informative of full problem
 - Easy to undershoot or overshoot the minimum

A Quick, Initial Comparison

Algorithm 1 ADAM

Input: $x_1 \in \mathbb{R}^d$, learning rate $\{\eta_t\}_{t=1}^T$, decay parameters $0 \leq \beta_1, \beta_2 \leq 1, \epsilon > 0$

Set $m_0 = 0, v_0 = 0$

for $t = 1$ **to** T **do**

 Draw a sample s_t from \mathbb{P} .

 Compute $g_t = \nabla \ell(x_t, s_t)$.

$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$

$v_t = v_{t-1} - (1 - \beta_2)(v_{t-1} - g_t^2)$

$x_{t+1} = x_t - \eta_t m_t / (\sqrt{v_t} + \epsilon)$

end for

Algorithm 2 YOGI

Input: $x_1 \in \mathbb{R}^d$, learning rate $\{\eta_t\}_{t=1}^T$, parameters $0 < \beta_1, \beta_2 < 1, \epsilon > 0$

Set $m_0 = 0, v_0 = 0$

for $t = 1$ **to** T **do**

 Draw a sample s_t from \mathbb{P} .

 Compute $g_t = \nabla \ell(x_t, s_t)$.

$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$

$v_t = v_{t-1} - (1 - \beta_2) \text{sign}(v_{t-1} - g_t^2) g_t^2$

$x_{t+1} = x_t - \eta_t m_t / (\sqrt{v_t} + \epsilon)$

end for

The Main Result: YOGI Method

Algorithm 2 YOGI

Input: $x_1 \in \mathbb{R}^d$, learning rate $\{\eta_t\}_{t=1}^T$, parameters

$0 < \beta_1, \beta_2 < 1, \epsilon > 0$

Set $m_0 = 0, v_0 = 0$

for $t = 1$ **to** T **do**

 Draw a sample s_t from \mathbb{P} .

 Compute $g_t = \nabla \ell(x_t, s_t)$.

$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$

$v_t = v_{t-1} - (1 - \beta_2) \text{sign}(v_{t-1} - g_t^2) g_t^2$

$x_{t+1} = x_t - \eta_t m_t / (\sqrt{v_t} + \epsilon)$

end for

Overview

- Magnitude depends only on square of gradient

Advantages

- More control over learning weights
- $O(d)$ memory, $O(d)$ time complexity, $O(1/\delta^2)$ SFO
- The sign term gives previous gradient information influence on the direction, but not size of the second moment term

Drawbacks

- Performance compared to ADAM is nearly negligible

In-Depth

(Note that this slide is not shown in the video presentation. We provide more intuition of the differences of the algorithms here, for the interested viewer.)

What makes the methods different?:

Suppose g_t is small, and v_{t-1} is large.

ADAM:

Then $(v_{t-1} - g_t^2)$ is also large, which means v_t is decreased by a fraction of a large amount. Then, our stepsize increases by a larger multiplicative factor of $\frac{1}{\sqrt{v_t}}$. However, we want small stepsize for small g_t .

YOGI:

Then $(v_{t-1} - g_t^2)$ is positive, but since g_t is small, we decrease v_t by a fraction of a small amount. Then, our stepsize increases by a smaller multiplicative factor of $\frac{1}{\sqrt{v_t}}$. That is, we increase the stepsize less compared to ADAM, when the gradient suddenly goes to 0, which is common in sparse settings.

There's also the case to consider when g_t is large, and v_{t-1} is slightly larger. In ADAM, v_t decreases slightly, and so the stepsize increases only slightly. In YOGI, v_t decreases, which means our step-size increases. YOGI here takes a more aggressive increase in step-size when it sees our gradients decrease.

Finally, when g_t is larger, and v_{t-1} is large, Then v_t increases slightly for ADAM, which means a reduction in step-size. In YOGI, we see v_t increases by a larger factor, which means a more aggressive reduction of step-size when we see the gradient increase.

Algorithm 1 ADAM

Input: $x_1 \in \mathbb{R}^d$, learning rate $\{\eta_t\}_{t=1}^T$, decay parameters $0 \leq \beta_1, \beta_2 \leq 1, \epsilon > 0$
Set $m_0 = 0, v_0 = 0$
for $t = 1$ **to** T **do**
 Draw a sample s_t from \mathbb{P} .
 Compute $g_t = \nabla \ell(x_t, s_t)$.
 $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
 $v_t = v_{t-1} - (1 - \beta_2)(v_{t-1} - g_t^2)$
 $x_{t+1} = x_t - \eta_t m_t / (\sqrt{v_t} + \epsilon)$
end for

Algorithm 2 YOGI

Input: $x_1 \in \mathbb{R}^d$, learning rate $\{\eta_t\}_{t=1}^T$, parameters $0 < \beta_1, \beta_2 < 1, \epsilon > 0$
Set $m_0 = 0, v_0 = 0$
for $t = 1$ **to** T **do**
 Draw a sample s_t from \mathbb{P} .
 Compute $g_t = \nabla \ell(x_t, s_t)$.
 $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
 $v_t = v_{t-1} - (1 - \beta_2) \text{sign}(v_{t-1} - g_t^2) g_t^2$
 $x_{t+1} = x_t - \eta_t m_t / (\sqrt{v_t} + \epsilon)$
end for

Takeaways Compared to ADAM:

Yogi increases step-size in a slow, controlled manner when the previous curvature estimate is large, and our current stochastic gradient is small. (Reduce overshooting in sparse settings)

Yogi increases step-size aggressively when our estimate is large, and our current stochastic gradient is slightly smaller. (Improve undershooting)

Yogi decreases step-size aggressively when our current stochastic gradient is slightly larger than our large curvature estimate. (Possibly reduce overshooting)

The sign() in Yogi keeps the “direction” of how our step-size changes normally in Adam: We want large recent gradients to decrease step-size in future iterations, and small recent gradients to increase step-size in future iterations.

Time and Memory Complexity

	SFO Complexity (Convergence - We assume $b = \Theta(T)$)	Memory Costs	Computational Cost per Iteration (mini-batch = 1)
SGD	$O(\frac{1}{\delta^2})$	$O(d)$	$O(d)$
ADAM	$O(\frac{1}{\delta^2})$	$O(d)$	$O(d)$
YOGI	$O(\frac{1}{\delta^2})$	$O(d)$	$O(d)$

- Equivalent convergence & complexity
- Computational Cost for mini-batch > 1 is $O(bd)$
- SFO Complexity of ADAM & YOGI with large mini-batch is equivalent to SGD

Performance Guarantees and Algorithm Analysis

Corollary 4. For x_t generated using YOGI with constant η (and parameters from Theorem 2), we have

$$\mathbb{E}[\|\nabla f(x_a)\|^2] \leq O\left(\frac{1}{T} + \frac{1}{b}\right)$$

where x_a is an iterate uniformly randomly chosen from $\{x_1, \dots, x_T\}$.

-Some assumptions: $1 - \beta_2 \leq \frac{\epsilon^2}{16G^2}$ and $\eta \leq \frac{\epsilon\sqrt{\beta_2}}{2L}$

Corollary 5. YOGI with $b = \Theta(T)$ and constant η (and parameters from Theorem 2) has SFO complexity is $O(1/\delta^2)$ for achieving a δ -accurate solution.

- Expected stationarity of the objective function is bounded by the number of max iterations T , and mini-batch size b .
- SFO Complexity of $O(1/\delta^2)$ is achieved when b is tightly bound to T , as T goes to infinity.


Interpretations: $\frac{1}{T} \sum_{t=1}^T \mathbb{E}[\|\nabla f(x_t)\|^2] \leq 2(\sqrt{2}G + \epsilon) \times \left[\frac{f(x_1) - f(x^*)}{T} + \left(\frac{G\sqrt{1-\beta_2}}{\epsilon^2} + \frac{L\eta}{2\epsilon^2\sqrt{\beta_2}} \right) \frac{\sigma^2}{b} \right]$

(G - Upper bound of gradient per element)
 (L - Lipschitz Constant, where the the loss f is L -smooth)
 (σ^2 - Upper bound on variance of stochastic gradients)

In Convex
 Optimization: $f(x) - f(x^*)$

Empirical Evaluations - Multilogistic Regression on Fashion MNIST

$\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 1e-8$, $\alpha = 0.1$ $\eta = 1.0/(1.0+\alpha t)$




Empirical Evaluations - Adaptive Methods Comparisons


$$\text{Objective Function } f = \frac{1}{2} * ||Dx||^2$$

$$\beta_1 = 0.9, \beta_2 = 0.999, \varepsilon = 1e-8, \alpha = 1.0, \eta = 1.0/(1.0+\alpha t)$$

Contours of f



Gradient Size Over Iterations




Empirical Evaluations - Autoencoder with Fashion MNIST


Architecture: Hidden Layer 1: ReLU
100 Neurons, Binary Cross Entropy Loss,
20 Iterations, minibatch = 25

Optimizer	Train Loss	Test Loss
SGD $\eta = 10^{-4}$ $a = 10^{-3}$	0.42403	0.42523
ADAM $\beta_1 = 0.9$ $\beta_2 = 0.999$ $\epsilon = 10^{-8}$	0.38119	0.38337
YOGI $\beta_1 = 0.9$ $\beta_2 = 0.999$ $\epsilon = 10^{-3}$	0.38158	0.38379

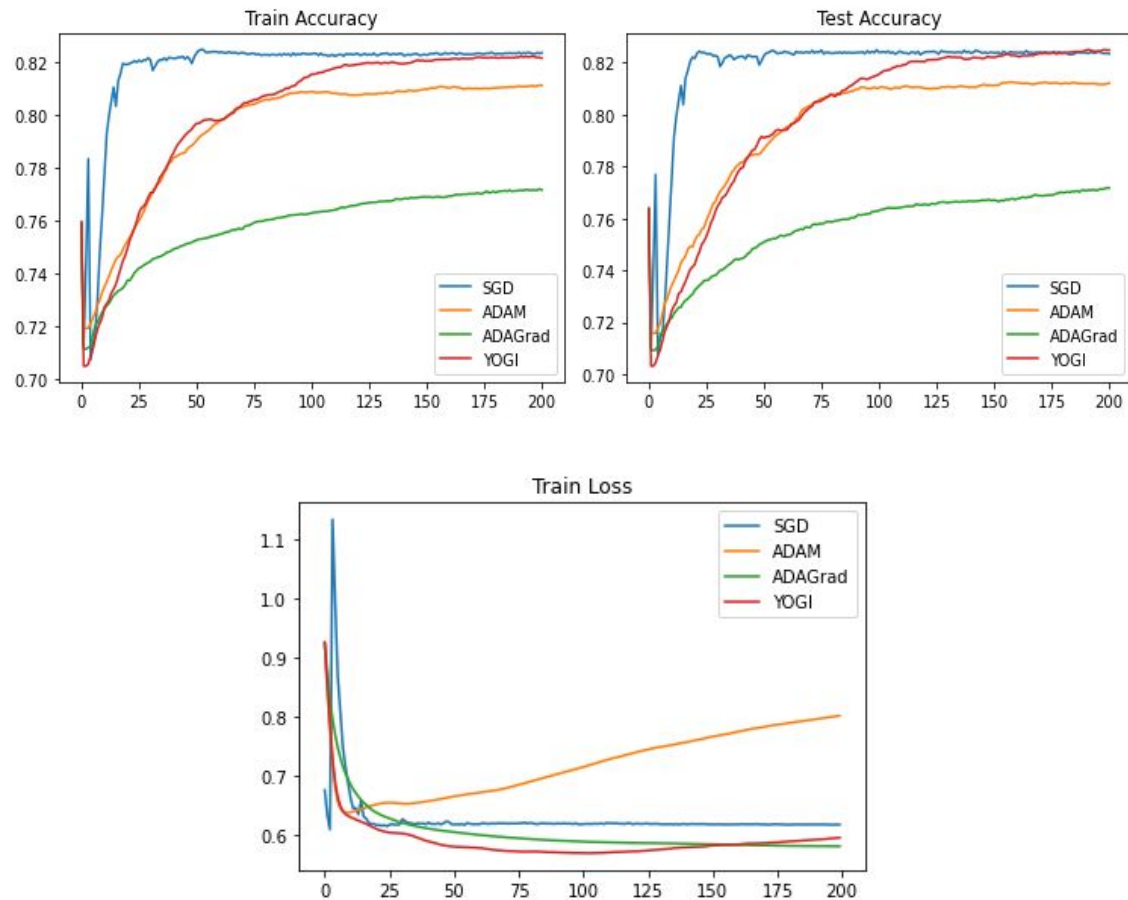
Loss Comparison



Test Loss



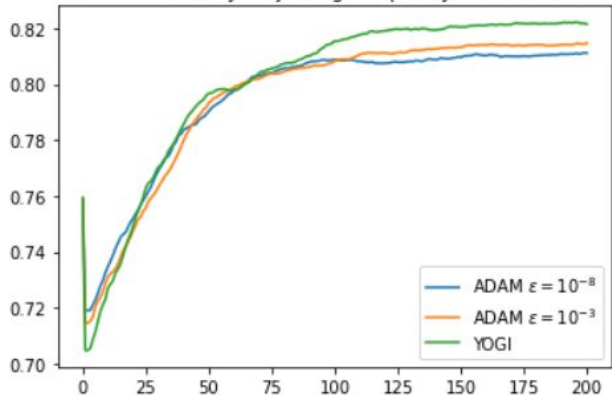
Empirical Evaluations - UCI Adult Classification SVM




General Parameters Used: Mini-batch = 1000	Test Accuracy	Train Accuracy	Train Loss
Methods Below:			
SGD $\alpha = \frac{\eta}{1+a \cdot t}$ $\eta = 0.1$	0.82341	0.82343	0.6178757
ADAGrad $\eta = 0.01$ $\epsilon = 10^{-8}$	0.77169	0.77150	0.581468
ADAM $\eta = 0.01$ $\beta_1 = 0.9$ $\beta_2 = 0.999$ $\epsilon = 10^{-8}$	0.81211	0.81112	0.80178
YOGI $\eta = 0.01$ $\beta_1 = 0.9$ $\beta_2 = 0.999$ $\epsilon = 10^{-3}$	0.82482	0.82147	0.59588

Empirical Evaluations - UCI Adult Classification SVM


Train Accuracy, Adjusting Adaptivity Parameter ϵ



Test Accuracy, Adjusting Adaptivity Parameter ϵ



Train Loss, Adjusting Adaptivity Parameter ϵ



General Parameters Used: Mini-batch = 1000 $\beta_1 = 0.9$ $\beta_2 = 0.999$ Methods Below:	Test Accuracy	Train Accuracy	Train Loss
ADAM $\epsilon = 10^{-8}$	0.81211	0.81112	0.80178
ADAM $\epsilon = 10^{-3}$	0.81039	0.81462	0.69412
YOGI $\epsilon = 10^{-3}$	0.82482	0.82147	0.59588

Algorithm 2 YOGI

Input: $x_1 \in \mathbb{R}^d$, learning rate $\{\eta_t\}_{t=1}^T$, parameters $0 < \beta_1, \beta_2 < 1, \epsilon > 0$
Set $m_0 = 0, v_0 = 0$
for $t = 1$ **to** T **do**
 Draw a sample s_t from \mathbb{P} .
 Compute $g_t = \nabla \ell(x_t, s_t)$.
 $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
 $v_t = v_{t-1} - (1 - \beta_2) \text{sign}(v_{t-1} - g_t^2) g_t^2$
 $x_{t+1} = x_t - \eta_t m_t / (\sqrt{v_t} + \epsilon)$
end for

Use Cases

- The performance of YOGI is comparable to ADAM
- If convergence is an issue, YOGI is a more robust choice than ADAM
- SGD can outperform YOGI/ADAM, but this requires either a highly detailed knowledge of your problem space or a large amount of time tuning your hyperparameters.

Problem Set

- One of the main results of YOGI is that it can be shown that the bound on the stationary condition decreases linearly with increased batch size
 - Implement YOGI in your HW6 autoencoder.
 - Run your autoencoder with YOGI with minibatch sizes of 16, 32, 64, 128
 - Comment on results
- The paper also stated that the optimal YOGI parameters are $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 1e-3$
 - With a minibatch size of 128 in your autoencoder, run YOGI with those parameters, and some others of your choice
 - Discuss your observations, and say whether you agree with the paper that those are the optimal parameters.

References & Supplements

Sashank J. Reddi, Manzil Zaheer, Devendra Sachan, Satyen Kale, Sanjiv Kumar. Adaptive Methods for Nonconvex Optimization. NeurIPS, 2018.

Sashank J Reddi, Satyen Kale, and Sanjiv Kumar. On the convergence of adam and beyond.arXiv preprint arXiv:1904.09237, 2019.

Dbouk Hassan. On The Convergence of SGD, ADAM & AMS-GRAD. ECE 543 Project Report, 2019.